

NCKU_Computer_Organization_2024_Assignment_1.md

µRISC-V: An Enhanced RISC-V Processor Design using Spike

Computer Organization 2024 Programming Assignment I

Due Date: 23:59, April 10, 2024

Overview

This assignment aims to get familiar with the hardware/software interface: instruction set architecture (ISA), as well as the environment and tools (e.g., compiler and RISC-V simulator) for RISC-V programming. In particular, this programming assignment is to **write inline assembly code** and run the developed code on a RISC-V ISA Simulator, *Spike*, to obtain the simulation results.

Prerequisite

If you have completed the environment setup based on the HW0 document, you can skip the following content and jump to 1. RISC-V Development Tools Installation.

The homework assignments of this course are based on the RISC-V Spike simulator running on Ubuntu Linux 22.04. As our assignments are built upon open source projects, there will be unexpected compatibility issues if you choose a different development environment.

For those students who use Mac or Windows systems, in order to save your time for establishing your development environment, our suggestion is the virtual machine based solution to work on your assignments. Please follow the procedures below to establish your *virtual* environment with Oracle VirtualBox. You can use any other virtual machine software (e.g., VMware's solution) to build your virtual environment.

- Install [Oracle VirtualBox](#)
 - Choose and download the Oracle VM VirtualBox Base Package based on the *operating system* of your host machine.
 - If you have a machine with Apple Silicon processors, including M1 and M2, and you cannot install the latest VirtualBox software on your machine, you may try to install an older version (VirtualBox-7.0.2_BETA4-154219-macOSArm64.dmg) from [VirtualBox download site](#).
- Download the pre-build Ubuntu 22.04 image for virtual machines.
 - Option #1: [Ubuntu 22.04 Linux VM Images for VirtualBox and VMware](#)
 - Be aware of the username and password for the pre-built system, and the HW resources allocated for the virtual machine.
 - Username: ubuntu; Password : ubuntu;
 - vCPU: 2, RAM: 4 GB, DISK: 128GB, Network: NAT
 - Option #2: You can find some other online resources to help you.
 - For example, Ubuntu provides a tutorial [How to run an Ubuntu Desktop virtual machine using VirtualBox 7](#)
- Create a virtual machine for Ubuntu 22.04 with VirtualBox.
 - You can refer to the [tutorial from Ubuntu](#) to create a virtual machine with Ubuntu 22.04.

- Or, you can find some other online resources to help you.

1. RISC-V Development Tools Installation




Figure 1. The execution flow for running C code on Spike.

The figure above illustrates the execution flow of a C program with Spike, a RISC-V processor simulator, on top of your host machine. There are three yellow boxes represents three different tools: *riscv-gnu-toolchain*, *riscv-isa-sim*, and *riscv-pk*.

The first software, *riscv-gnu-toolchain*, is the RISC-V C and C++ cross-compiler. It supports two build modes: a generic ELF/Newlib toolchain and a more sophisticated Linux-ELF/glibc toolchain. The instructions to install this software on your virtual environment are summarized in Section 1-1. You can also refer to its open source repository, [RISC-V GNU Compiler Toolchain](#), for more information.

The second software, *riscv-isa-sim*, is a RISC-V simulator called *Spike* that implements a functional model of a variety of RISC-V designs. It is named after the golden spike used to celebrate the completion of the US transcontinental railway. The installation instructions are given in Section 1-2. You can find out the detailed information (e.g., the supported RISC-V extensions) in its open source repository: [riscv-isa-sim](#).

The third software, *riscv-pk*, is a lightweight application execution environment that can host statically-linked RISC-V ELF binaries. It is designed to support tethered RISC-V implementations with limited I/O capability and thus handles I/O-related system calls by proxying them to a host computer. In order words, given a C program, the ordinary RISC-V instructions are emulated by Spike and the system calls (system services) are served by your virtual environment (i.e., the Ubuntu 22.04), where the system calls are forwarded by *riscv-pk* so that the underlying OS can take care of the system calls invoked by your C program. For more information, you can refer to its open source repository, [riscv-pk](#).

1-1. Install RISC-V GNU Compiler Toolchain

This is the RISC-V C and C++ cross-compiler. You should find a location to install the cross-compiler. You can create a package folder under the `/opt` folder (i.e., `/opt/riscv`) and add this into the environment variable `$RISCV`. The following example assumes the package is installed under the path: `/opt/riscv`. If you choose another path for the installation, you should replace the string of `/opt/riscv` with your selected path in the following commands.

After you choose the folder for the package, you should also create a subfolder `bin` within your package folder (e.g., `/opt/riscv/bin`) and add the `bin` folder into to your system PATH. The related commands are listed below.

```

$ sudo mkdir /opt/riscv && sudo mkdir /opt/riscv/bin
$ export RISCV=/opt/riscv >> ~/.bashrc
$ export PATH="$PATH": "$RISCV"/bin >> ~/.bashrc
$ source ~/.bashrc

```

Later, you should create a project folder (e.g., `$HOME/riscv`). You should clone the above three tools and write your homework in project folder. The directory tree is shown as below.

```
$HOME/riscv/  
├── riscv-gnu-toolchain  
├── riscv-pk  
├── riscv-isa-sim  
└── CO_StudentID_HW1/
```

The following commands are used to clone the riscv-gnu-toolchain project and installing related software packages. Note that `./configure --prefix=...` is used to configure the prefix (the path of the package folder); the prefix is assumed to be: `/opt/riscv`. In addition the last command `sudo make -j4` is used to build the project from source files with four threads in parallel.

```
$ mkdir ~/riscv  
$ cd ~/riscv  
$ git clone https://github.com/riscv/riscv-gnu-toolchain  
$ sudo apt-get install autoconf automake autotools-dev curl python3 libmpc-dev  
libmpfr-dev libgmp-dev gawk build-essential bison flex texinfo gperf libtool  
patchutils bc zlib1g-dev libexpat-dev  
$ cd riscv-gnu-toolchain  
$ ./configure --prefix=/opt/riscv --enable-multilib  
$ sudo make linux -j4
```

After you do the above operations, riscv-gnu-toolchain will be installed under the path: `$HOME/riscv/riscv-gnu-toolchain` and riscv-gnu-toolchain packages will be installed under the path: `/opt/riscv/riscv64-unknown-linux-gnu`.

1-2. Install RISC-V ISA Simulator

Here, we start to install the RISC-V ISA simulator. Please go to the project directory (e.g., `/$HOME/riscv`) and run the following commands to clone the riscv-isa-sim project.

```
$ cd ~/riscv  
$ sudo apt install device-tree-compiler  
$ git clone https://github.com/riscv-software-src/riscv-isa-sim.git  
$ cd riscv-isa-sim  
$ mkdir build  
$ cd build  
$ ../configure --prefix=/opt/riscv  
$ make  
$ sudo make install
```

After you do the above operations, riscv-isa-sim will be installed under the path: `$HOME/riscv/riscv-isa-sim`.

1-3. Install RISC-V Proxy Kernel

We start to install the proxy kernel software. Please go to the project directory (e.g., `$HOME/riscv`) and run the following commands to clone the `riscv-pk` project.

```
$ cd ~/riscv
$ git clone https://github.com/riscv-software-src/riscv-pk.git
$ cd riscv-pk
$ mkdir build
$ cd build
$ ../configure --prefix=/opt/riscv --host=riscv64-unknown-linux-gnu
$ make
$ sudo make install
```

After you do the above operations, `riscv-isa-sim` will be installed under the path: `$HOME/riscv/riscv-pk`.

1-4. Run a simple code on Spike

To validate the above tools are installed correctly, you can compile (with the compiler `riscv64-unknown-linux-gnu`) and run an example program (e.g., `hello.c`) with Spike. In the current version, Spike only supports the statically built machine executable, so we need to compile the example program with the flag `-static`.

The second command suggests that the Spike simulator supports the `RV64GC` RISC-V variant (More about the RISC-V ISA base and extensions can refer to [the page](#)). Also, the binary executable `hello` is run with the support of the built proxy kernel software under the path: `/opt/riscv/riscv64-unknown-linux-gnu/bin/pk`.

```
$ riscv64-unknown-linux-gnu-gcc -static -o hello hello.c
$ spike --isa=RV64GC /opt/riscv/riscv64-unknown-linux-gnu/bin/pk hello
```

If the above two commands are performed successfully, the terminal will show the following contents.

```
Hello
```

2. GCC Inline Assembly

The inline assembly code provide a way to write efficient code. One of the benefits of the inline assembly is the reducing of the overheads incurred by function calls. The RISC-V C/C++ compiler is based on GCC compilers, and GCC inline assembly uses AT&T/UNIX assembly syntax. The basic format of an inline assembly code is define as below. You can see the following code example to get a high-level concept of [the format of inline assembly](#). For detailed information, you would refer to [the extended asm. from GCC HOWTO document](#).

```
asm volatile( AssemblerTemplate
              : OutputOperands
              : InputOperands
              : Clobbered registers)
```

Example:

Original C program: `add.c`

```
#include <stdio.h>

int main ()
{
    int a = 10, b = 5;
    a = a + b;
    printf("%d\n", a);
    return 0;
}
```

Inline assembly version: `add_inline.c`

```
#include <stdio.h>

int main ()
{
    int a = 10, b = 5;
    //a = a + b;
    asm volatile(
        "add %[a], %[a], %[b]\n\t" //AssemblerTemplate
        :[a] "+r"(a)              //OutputOperands, "=r" means write-only, "+r"
means read/write
        :[b] "r"(b)                //InputOperands
    );
    printf("%d\n", a);
    return 0;
}
```

- "a" is the output operand, referred to by the register %[a] and "b" is the input operand, referred to by the register %[b]. In inline assembly, you can adopt the `asmSymbolicName` syntax method which is to renaming the registers (e.g., [a], [b]). The benefit for the `asmSymbolicName` syntax method is more readable and more maintainable since reordering index numbers is not necessary when adding or removing operands.
- "r" is a constraint on the operands. "r" says to GCC to use any register for storing the operands. The constraint modifier "=" says the output operand is write-only. The constraint modifier "+" says the output operand can both read and write.
- Beware! Every line of instruction should be ended with `\n\t`.

3. Assignment

This assignment has four exercises, 3-1~3-4, a total of 100 pt. Please write your code in `asm volatile()`. Any modifications outside of `asm volatile()` are not allowed.

For the exercises, 3-1 and 3-2, the students are divided into three groups according each student's Student ID. You are responsible for complete the code which according to your group.

- Group 1: StudentID mod 3 = 1.
- Group 2: StudentID mod 3 = 2.
- Group 3: StudentID mod 3 = 0.

3-1 Basic (20%)

Please finish the following code which according to your group. You need to add your inline assembly code in the C file, belonging to your group.

- Group 1: sub_inline.c

```
#include <stdio.h>

int main ()
{
    int a, b;
    FILE *input = fopen("../input/1.txt", "r");
    fscanf(input, "%d %d", &a, &b);
    fclose(input);
    //a = a - b;
    asm volatile(/*Your Code*/);
    printf("%d\n", a);
    return 0;
}
```

- Group 2: mul_inline.c

```
#include <stdio.h>

int main ()
{
    int a, b;
    FILE *input = fopen("../input/1.txt", "r");
    fscanf(input, "%d %d", &a, &b);
    fclose(input);
    //a = a * b;
    asm volatile(/*Your Code*/);
    printf("%d\n", a);
    return 0;
}
```

- Group 3: div_inline.c

```
#include <stdio.h>

int main ()
{
    int a, b;
    FILE *input = fopen("../input/1.txt", "r");
    fscanf(input, "%d %d", &a, &b);
    fclose(input);
    //a = a / b;
    asm volatile(/*Your Code*/);
    printf("%d\n", a);
    return 0;
}
```

3-2 Array (20%)

Please finish the following code which according to your group. The original C code is in the comments.

- Group 1: subarray_inline.c

```
#include <stdio.h>

int main ()
{
    int a[10] = {0}, b[10] = {0}, c[10] = {0};
    int i, arr_size = 10;
    FILE *input = fopen("../input/2.txt", "r");
    for(i = 0; i < arr_size; i++) fscanf(input, "%d", &a[i]);
    for(i = 0; i < arr_size; i++) fscanf(input, "%d", &b[i]);
    for(i = 0; i < arr_size; i++) fscanf(input, "%d", &c[i]);
    fclose(input);
    int *p_a = &a[0];
    int *p_b = &b[0];
    int *p_c = &c[0];

    /* original code
    for (int i = 0; i < arr_size; i++){
        p_c[i] = p_a[i] - p_b[i];
    }
    */
    for (int i = 0; i < arr_size; i++){
        asm volatile(/*Your Code*/);
    }
    p_c = &c[0];
    for(i = 0; i < arr_size; i++) printf("%d ", *p_c++);
    printf("\n")
    return 0;
}
```

- Group 2: mularray_inline.c

```
#include <stdio.h>

int main ()
{
    int a[10] = {0}, b[10] = {0}, c[10] = {0};
    int i, arr_size = 10;
    FILE *input = fopen("../input/2.txt", "r");
    for(i = 0; i < arr_size; i++) fscanf(input, "%d", &a[i]);
    for(i = 0; i < arr_size; i++) fscanf(input, "%d", &b[i]);
    for(i = 0; i < arr_size; i++) fscanf(input, "%d", &c[i]);
    fclose(input);
    int *p_a = &a[0];
    int *p_b = &b[0];
    int *p_c = &c[0];

    /* original code
    for (int i = 0; i < arr_size; i++){
        p_c[i] = p_a[i] * p_b[i];
    }
    */
    for (int i = 0; i < arr_size; i++){
        asm volatile(/*Your Code*/);
    }
    p_c = &c[0];
    for(i = 0; i < arr_size; i++) printf("%d ", *p_c++);
    printf("\n")
    return 0;
}
```

- Group 3: divarray_inline.c

```
#include <stdio.h>

int main ()
{
    int a[10] = {0}, b[10] = {0}, c[10] = {0};
    int i, arr_size = 10;
    FILE *input = fopen("../input/2.txt", "r");
    for(i = 0; i < arr_size; i++) fscanf(input, "%d", &a[i]);
    for(i = 0; i < arr_size; i++) fscanf(input, "%d", &b[i]);
    for(i = 0; i < arr_size; i++) fscanf(input, "%d", &c[i]);
    fclose(input);
    int *p_a = &a[0];
    int *p_b = &b[0];
    int *p_c = &c[0];

    /* original code
```



```

    for (int i = 0; i < arr_size; i++){
        p_c[i] = p_a[i] / p_b[i];
    }
    /*
    for (int i = 0; i < arr_size; i++){
        asm volatile(/*Your Code*/);
    }
    p_c = &c[0];
    for(i = 0; i < arr_size; i++) printf("%d ", *p_c++);
    printf("\n")
    return 0;
}

```

3-3 Matrix Multiply (30%)

Given the matrix multiplication code `matrix.c` below, please convert the **loop body** of the nested for-loop into the corresponding inline assembly version.

Original C Code:

```

/*
 * description:          matrix - multiply benchmarking
 *
 * |h11 h12 h13| |x1 x2| |y1 y2| | h11*x1+h12*x3+h13*x5 h11*x2+h12*x4+h13*x6|
 * |h21 h22 h23|*|x3 x4|=|y3 y4|=| h21*x1+h22*x3+h23*x5 h21*x2+h22*x4+h23*x6|
 * |h31 h32 h33| |x5 x6| |y5 y6| | h31*x1+h32*x5+h33*x5 h31*x2+h32*x4+h33*x6|
 *
 * Element are to store in following order:
 *
 * matrix h[9]={h11,h12,h13, h21,h22,h23, h31,h32,h33}
 * vector x[6]={x1,x2, x3,x4, x5,x6}
 * vector y[6]={y1,y2, y3,y4, y5,y6}
 */

#include<stdio.h>

int main()
{
    int f,i,j;
    int h[9] = {0}, x[6] = {0}, y[6] = {0};
    for(i = 0; i < 9; i++) scanf("%d", &h[i]);
    for(i = 0; i < 6; i++) scanf("%d", &x[i]);
    for(i = 0; i < 6; i++) scanf("%d", &y[i]);

    int *p_x = &x[0];
    int *p_h = &h[0];
    int *p_y = &y[0];

    for (i = 0; i < 3; i++){
        /* p_x points to the beginning of the input vector */
        p_x = &x[0];

```

```

    /* do matrix multiply */
    for (j = 0; j < 2; j++){
        p_h = &h[i * 3];
        for (f = 0; f < 3; f++)
            *p_y += *p_h++ * *p_x++;
        *p_x++;

        /* next element */
        p_x = &x[1];
        p_y++;
    }
}
return 0;
}

```

Please modify the above c code to inline assembly version. You should copy the following code and write your code in `asm volatile()`. Beware Your Code is **in** the for loop.

```

#include<stdio.h>

int main()
{
    int f, i, j;
    int h[9] = {0}, x[6] = {0}, y[6] = {0};
    FILE *input = fopen("../input/3.txt", "r");
    for(i = 0; i < 9; i++) fscanf(input, "%d", &h[i]);
    for(i = 0; i < 6; i++) fscanf(input, "%d", &x[i]);
    for(i = 0; i < 6; i++) fscanf(input, "%d", &y[i]);
    fclose(input);

    int *p_x = &x[0];
    int *p_h = &h[0];
    int *p_y = &y[0];

    for (i = 0; i < 3; i++){
        for (j = 0; j < 2; j++){
            for (f = 0; f < 3; f++){
                asm volatile(/*Your Code*/);
            }
        }
    }
    p_y = &y[0];
    for(i = 0; i < 6; i++) printf("%d ", *p_y++);
    printf("\n");
    return 0;
}

```

3-4 Matrix Multiply (30%)

Given the matrix multiplication code `matrix.c` in 3-3, please convert the nested for-loop into the inline assembly version. You should copy the following code and write your code in `asm volatile()`. Note that the entire nested for-loop should be converted into the inline assembly code in this exercise.

```
#include<stdio.h>

int main()
{
    int f, i = 0;
    int h[9] = {0}, x[6] = {0}, y[6] = {0};
    FILE *input = fopen("../input/4.txt", "r");
    for(i = 0; i < 9; i++) fscanf(input, "%d", &h[i]);
    for(i = 0; i < 6; i++) fscanf(input, "%d", &x[i]);
    for(i = 0; i < 6; i++) fscanf(input, "%d", &y[i]);
    fclose(input);

    int *p_x = &x[0];
    int *p_h = &h[0];
    int *p_y = &y[0];

    asm volatile(/*Your Code*/);

    p_y = &y[0];
    for(i = 0; i < 6; i++) printf("%d ", *p_y++);
    printf("\n");
    return 0;
}
```

4. Test Your Assignment

We use `local-judge` to judge your program. You can use the command `pip3 install local-judge` to download and install the local-judge, as shown below.

```
$ sudo apt install python3-pip # Install pip tool
$ pip3 install local-judge    # Install local-judge via pip3
```

The source files of the aforementioned exercises are included in the zip file `C02024HW1.zip` that can be downloaded from the course website on NCKU Moodle. Please unzip and move these files into the project folder(e.g., `$HOME/riscv`). The file `judge*.conf` should be placed in the same folder as your code(e.g., `$HOME/riscv/C0_StudentID_HW1`).The directory tree of the RISC-V tools and the source code of this assignment is shown as below.

```
$HOME/riscv/
├── riscv-gnu-toolchain
├── riscv-pk
└── riscv-isa-sim
```

```

├── CO_StudentID_HW1/
│   ├── 1.c
│   ├── 2.c
│   ├── 3.c
│   ├── 4.c
│   ├── judge1.conf
│   ├── judge2.conf
│   ├── judge3.conf
│   ├── judge4.conf
│   └── Makefile
├── input
└── answer

```

Now, you can use the judge program to get the score of your developed code with the following commands.

```

$ cd ~/riscv/CO_StudentID_HW1
$ make judge
=====+=====
Sample | Accept
=====+=====
    1 | ✓
=====+=====
    2 | ✓
=====+=====
    3 | ✓
=====+=====
    4 | ✗
=====+=====
Obtained/Total scores: 70/100

```

If you want to judge a specific exercise, you can use the command `judge -c [CONFIG]`. For example, you can check the result of the first exercise 3-1 via the command `judge -c judge1.conf`.

```

$ cd ~/riscv/CO_StudentID_HW1
$ judge -c judge1.conf
=====+=====
Sample | Accept
=====+=====
    1 | ✓
=====+=====
Correct/Total problems: 1/1
Obtained/Total scores: 20/20

```

If your code generates wrong output, you can use the parameter `-v 1` while running the judge command to check the differences between your output and the correct answer. The example message is shown as below.

Example:
 standard answer is 5.
 your assignment output is 15.

```
$ cd ~/riscv/CO_StudentID_HW1
$ judge -c judge1.conf -v 1
=====
Sample | Accept
=====
      1 | X
=====
diff --git a/home/ubuntu/riscv/answer/1.out b/./output/1local_1709739995.out
index 7ed6ff8..60d3b2f 100755
--- a/home/ubuntu/riscv/answer/1.out
+++ b/./output/1local_1709739995.out
@@ -1,1 +1,1 @@
515

=====
Correct/Total problems: 0/1
Obtained/Total scores: 0/20
```

5. Submission of Your Assignment

- Compress your source files into a **zip** file.
- Submit your homework to NCKU Moodle.
- The file organization of your zip file should be as follows.
 - **NOTE: Change all CO_StudentID to your student ID number, e.g., F12345678_HW1.zip.** The example zip file for the student with the student ID, F12345678, is shown below.

```
F12345678_HW1.zip/
├── F12345678_HW1/
│   ├── 1.c
│   ├── 2.c
│   ├── 3.c
│   ├── 4.c
│   └── Makefile
```

!!! Incorrect format (either the file structure or file name) will lose 10 points. !!!

6. Reference

- [GCC-Inline-Assembly-HOWTO](#)
- [Extended Asm - Assembler Instructions with C Expression Operands](#)
- [RISC-V Assembly Programmer's Manual](#)
- [The RISC-V Instruction Set Manual Version 2.2](#)

